

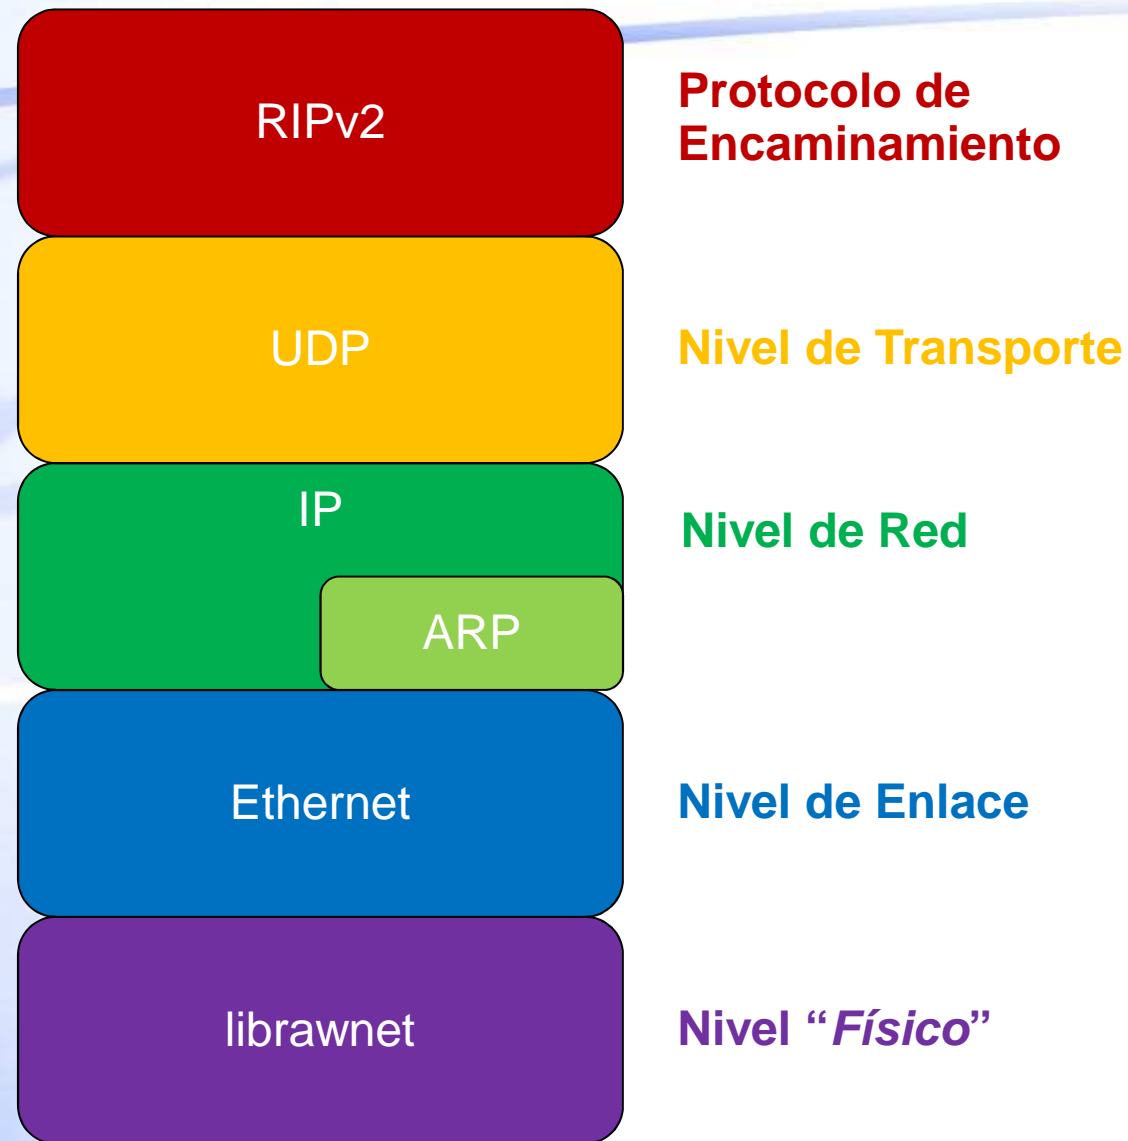
# Introducción a la librería de desarrollo librawnet

**Manuel Urueña <muruenya@it.uc3m.es>**

**Carlos J. Bernardos <cjbc@it.uc3m.es>**



# Pila de Protocolos RIPv2



# librawnet

## ◆ Librería para la programación de protocolos de red “*en crudo*”:


- ❖ Permite acceder directamente a los interfaces de red del equipo.

## ◆ Los paquetes enviados/recibidos no han sido procesados por la pila de protocolos del equipo:

- ❖ Es necesario añadir/procesar las **cabeceras** de todos los protocolos de la pila que se desea desarrollar.
- ❖ Realmente el *kernel* y librawnet reciben una **copia** de todos los paquetes recibidos por la tarjeta de red.

## ◆ librawnet proporciona dos módulos:

- ❖ Gestión de los interfaces de red, envío/recepción de paquetes “*crudos*”: `rawnet.h` (`man rawnet`)
- ❖ Temporizadores con precisión de milisegundos: `timerms.h` (`man timerms`)



# timerms.h

## ◆ Obtener tiempo actual:

```
long long int timerms_time();
```

## ◆ (Re)iniciar temporizador:

```
long long int timerms_reset ( timerms_t * timer,  
                             long int timeout );  
  
/* 'timeout' >= 0: 'timer' expira en 'timeout' ms  
   'timeout' < 0: Temporizador "infinito" */
```

## ◆ Consultar estado temporizador:

```
long int timerms_elapsed ( timerms_t * timer );  
long int timerms_left ( timerms_t * timer );
```



# timerms.h: Ejemplo de uso (I)

## ◆ Temporizador que expira a los 500 ms:

```
timerms_t timer;
long int timeout = 500;
long long int now = timerms_reset(&timer, timeout);
long elapsed_time = timerms_elapsed(&timer); /* elapsed_time == 0 */
long time_left = timerms_left(&timer);      /* time_left == 500 */
...
elapsed_time = timerms_elapsed(&timer);    /* elapsed_time == 100 */
time_left = timerms_left(&timer);          /* time_left == 400 */
...
elapsed_time = timerms_elapsed(&timer);    /* elapsed_time == 500 */
time_left = timerms_left(&timer);          /* time_left == 0 */
...
elapsed_time = timerms_elapsed(&timer);    /* elapsed_time == 1000 */
time_left = timerms_left(&timer);          /* time_left == 0 */
```



# timerms.h: Ejemplo de uso (II)

## ◆ Temporizador “*infinito*”:

```
timerms_t timer;  
long int timeout = -1;  
long long int now = timerms_reset(&timer, timeout);  
long elapsed_time = timerms_elapsed(&timer); /* elapsed_time == 0 */  
long int time_left = timerms_left(&timer);  /* time_left < -1 */  
...  
elapsed_time = timerms_elapsed(&timer);  /* elapsed_time == 1000 */  
time_left = timerms_left(&timer);         /* time_left < -1 */
```



# rawnet.h

## ◆ Operaciones con interfaces “*crudos*”:

```
rawiface_t * rawiface_open ( char* ifname );  
char* rawiface_getname ( rawiface_t * iface );  
int rawiface_getaddr ( rawiface_t * iface, char addr[] );  
int rawiface_getmtu ( rawiface_t * iface );  
int rawiface_close ( rawiface_t * iface );
```

## ◆ Enviar/Recibir paquetes “*crudos*”:


```
int rawnet_send ( rawiface_t * iface, char * pkt, int pkt_len );  
int rawnet_recv ( rawiface_t * iface, char buffer[], int buf_len,  
                 long int timeout );  
int rawnet_poll ( rawiface_t * ifaces[], int ifnum,  
                 long int timeout );
```

## ◆ Gestión de errores:

```
char* rawnet_strerror();
```



# Ejemplo: Cliente / Servidor Ethernet





## ◆ Direcciones MAC:

## eth.h

```
#define MAC_ADDR_SIZE 6

typedef unsigned char mac_addr_t [MAC_ADDR_SIZE];

mac_addr_t BROADCAST_MAC_ADDR = {0xFF, 0xFF, 0xFF, 0xFF, 0xFF, 0xFF};

#define MAC_ADDR_STR_LENGTH 18

void eth_mac_str ( mac_addr_t addr, char str[] );

int eth_str_mac ( char* str, mac_addr_t addr );
```

## ◆ Operaciones con interfaces Ethernet:

```
eth_iface_t * eth_open ( char* ifname );

char * eth_getname ( eth_iface_t * iface );

void eth_getaddr ( eth_iface_t * iface, mac_addr_t addr );

int eth_close ( eth_iface_t * iface );
```

## ◆ Enviar/Recibir tramas Ethernet:

```
int eth_send ( eth_iface_t * iface, mac_addr_t dst, uint16_t type,
               unsigned char * payload, int payload_len );

int eth_recv ( eth_iface_t * iface, mac_addr_t src, uint16_t type,
               unsigned char buffer[], long int timeout );

int eth_poll ( eth_iface_t * ifaces[], int ifnum, long int timeout );
```



# eth\_open( )

```
#include "eth.h"
#include <rawnet.h>
#include <timerms.h>

typedef struct eth_iface {
    rawiface_t * raw_iface; /* Manejador del interfaz "crudo" */
    mac_addr_t mac_address; /* Dirección MAC del interfaz. */
} eth_iface_t;

eth_iface_t * eth_open ( char* ifname ) {
    struct eth_iface * eth_iface;
    eth_iface = malloc(sizeof(struct eth_iface));

    rawiface_t * raw_iface = rawiface_open(ifname);
    if (raw_iface == NULL) { return NULL; }
    eth_iface->raw_iface = raw_iface;

    rawiface_getaddr(raw_iface, eth_iface->mac_address);
    return eth_iface;
}
```



# eth\_send( ) (I)

```
struct eth_frame {
    mac_addr_t dest_addr;           /* Dirección MAC destino */
    mac_addr_t src_addr;            /* Dirección MAC origen */
    uint16_t type;                  /* Campo 'Tipo' */
    unsigned char payload[ETH_MTU]; /* Campo 'Payload' */
};

int eth_send ( eth_iface_t * iface, mac_addr_t dst, uint16_t type,
               unsigned char * payload, int payload_len ) {

    /* Crear la Trama Ethernet y rellenar todos los campos */
    struct eth_frame eth_frame;
    memcpy(eth_frame.dest_addr, dst, MAC_ADDR_SIZE);
    memcpy(eth_frame.src_addr, iface->mac_address, MAC_ADDR_SIZE);
    eth_frame.type = htons(type);
    memcpy(eth_frame.payload, payload, payload_len);
    int eth_frame_len = ETH_HEADER_SIZE + payload_len;
```



# eth\_send( ) (II)

```
/* Imprimir trama Ethernet */
char* iface_name = eth_getname(iface);
eth_mac_str(dst, mac_addr_str);
printf("eth_send(type=0x%04x, payload[%d]) > %s/%s\n",
       type, payload_len, iface_name, mac_addr_str);
print_pkt((unsigned char *) &eth_frame, eth_frame_len, ETH_HEADER_SIZE);

/* Enviar la trama Ethernet creada con rawnet_send() */
bytes_sent = rawnet_send(iface->raw_iface, (unsigned char *)
                        &eth_frame, eth_frame_len);

if (bytes_sent == -1) {
    fprintf(stderr, "eth_send(): ERROR en rawnet_send(): %s\n",
            rawnet_strerror());
    return -1;
}
return (bytes_sent - ETH_HEADER_SIZE);
}
```



# eth\_rcv( ) (I)

```
int eth_rcv ( eth_iface_t * iface, mac_addr_t src, uint16_t type,
              unsigned char buffer[], long int timeout ) {
    timerms_reset(&timer, timeout);
    do {
        long int time_left = timerms_left(&timer);
        /* Recibir trama del interfaz Ethernet */
        unsigned char eth_buffer[ETH_FRAME_MAX_SIZE];
        int frame_len = rawnet_rcv (iface->raw_iface, eth_buffer,
                                    ETH_FRAME_MAX_SIZE, time_left);
        if ((frame_len == -1) || (frame_len == 0)) { return frame_len; }
        /* Comprobar si es la trama que nos han pedido */
        struct eth_frame * eth_frame_ptr = (struct eth_frame *) eth_buffer;
        int is_my_mac = (memcmp(eth_frame_ptr->dest_addr, iface->mac_address,
                                MAC_ADDR_SIZE) == 0);
        int is_target_type = (ntohs(eth_frame_ptr->type) == type);
    } while ( ! (is_my_mac && is_target_type) );
    ...
}
```



# eth\_rcv( ) (II)

```
...

/* Trama recibida con el 'Type' solicitado. */

/* Copiar datos y dirección MAC origen */
memcpy(src, eth_frame_ptr->src_addr, MAC_ADDR_SIZE);
payload_len = frame_len - ETH_HEADER_SIZE;
memcpy(buffer, eth_frame_ptr->payload, payload_len);

/* Imprimir trama recibida */
eth_mac_str(src, mac_addr_str);
char* iface_name = eth_getname(iface);
printf("eth_rcv(type=0x%04x, payload[%d]) < %s/%s\n",
       type, payload_len, iface_name, mac_addr_str);
print_pkt((unsigned char *) eth_frame_ptr, frame_len, ETH_HEADER_SIZE);

return payload_len;
}
```



# eth\_client.c

```
int main ( int argc, char * argv[] ){  
    /* Procesar argumentos */  
    if (argc != 5) {  
        printf("Uso: %s <iface> <tipo> <mac> <long>\n", argv[0]); exit(-1);  
    }  
    /* Abrir la interfaz Ethernet */  
    eth_iface_t * eth_iface = eth_open(iface_name);  
    if (eth_iface == NULL) { exit(-1); }  
    /* Enviar trama Ethernet al Servidor */  
    int err = eth_send(eth_iface, server_addr, eth_type, payload, payload_len);  
    if (err == -1) { exit(-1); }  
    /* Recibir trama Ethernet del Servidor y procesar errores */  
    len = eth_recv(eth_iface, src_addr, eth_type, buffer, 2000);  
    if (len <= 0) { fprintf(stderr, "%s: ERROR en eth_recv()\n", myself);  
    } else { printf("Recibidos %d bytes del Servidor Ethernet\n", len); }  
    /* Cerrar interfaz Ethernet */  
    eth_close(eth_iface);  
    return 0;  
}
```



# eth\_server.c

```
int main ( int argc, char * argv[] ){  
    /* Procesar parámetros */  
    if (argc != 3) { printf("Uso: %s <iface> <tipo>\n", argv[0]); exit(-1); }  
    /* Abrir la interfaz Ethernet */  
    eth_iface_t * eth_iface = eth_open(iface_name);  
    if (eth_iface == NULL) { exit(-1); }  
    while(1) {  
        /* Recibir trama Ethernet del Cliente */  
        printf("Escuchando tramas Ethernet (tipo=0x%04x) ...\n", eth_type);  
        int payload_len = eth_recv(eth_iface, src_addr, eth_type, buffer, timeout);  
        if (payload_len == -1) { break; }  
        /* Enviar la misma trama Ethernet de vuelta al Cliente */  
        int len = eth_send(eth_iface, src_addr, eth_type, buffer, payload_len);  
        if (len == -1) { break; }  
        printf("Enviado %d bytes al Cliente Ethernet:\n", payload_len);  
    }  
    /* Cerrar interfaz Ethernet */  
    eth_close(eth_iface);  
    return 0;  
}
```





# Desarrollar protocolos con librawnet

## ◆ Acceder a los interfaces y enviar/recibir paquetes en crudo requiere permisos especiales

❖ **Compilar con:** `rawnetcc /tmp/<prog> <f1.c> <f2.c> ...`

## ◆ **Ejemplo:** `rawnetcc /tmp/eth_client eth.c eth_client.c`

### 1. **Compila el ejecutable:**

```
/usr/bin/gcc -o /tmp/eth_client eth.c eth_client.c -lrawnet
```

### 2. **Otorga permisos especiales:**

```
/usr/sbin/setcap cap_net_raw,cap_net_admin=eip /tmp/eth_client
```

ó

```
/usr/bin/chown root.root /tmp/eth_client
```

```
/usr/bin/chmod 4755 /tmp/eth_client
```

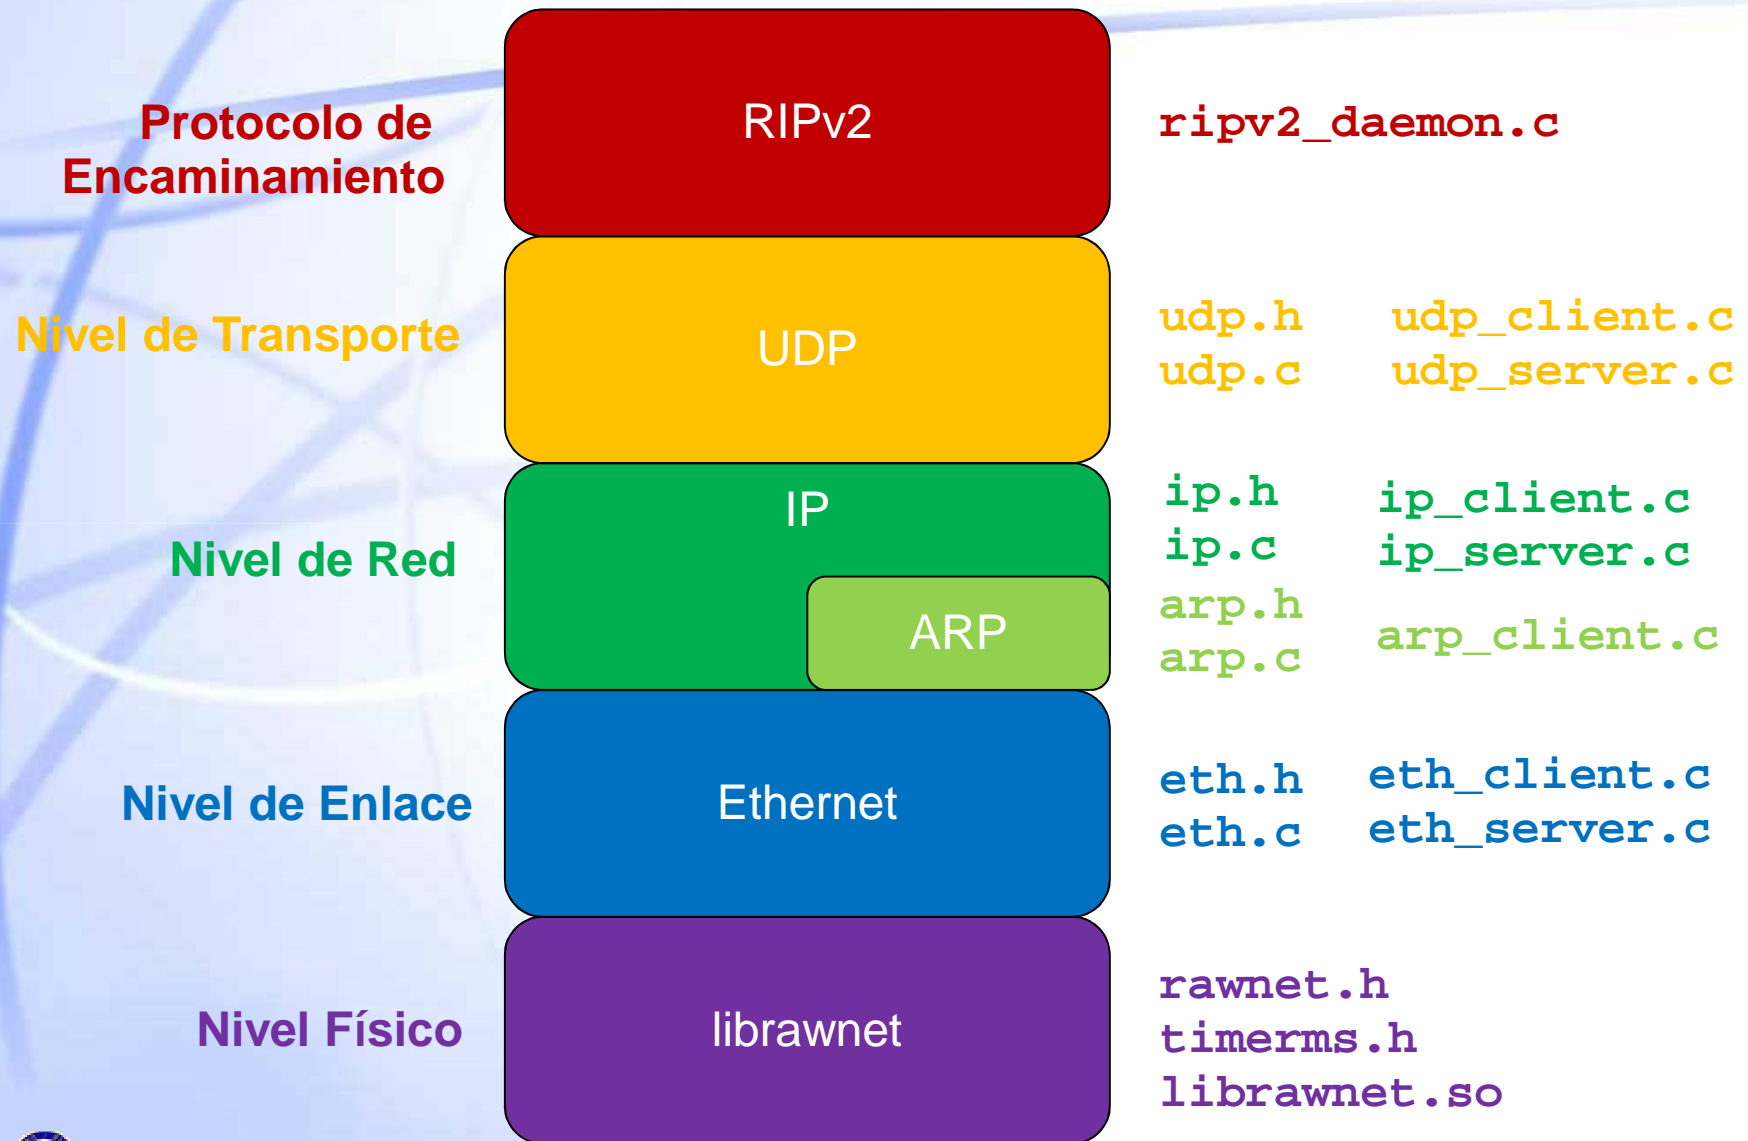
## ◆ **Depurar tráfico enviado con Wireshark**

❖ `xhost +`

❖ `sudo wireshark`



# Pila de Protocolos RIPv2



# arp.h

```
#include "eth.h"
```

```
#define IPv4_ADDR_SIZE 4
```

```
typedef unsigned char ipv4_addr_t[IPv4_ADDR_SIZE];
```

```
int arp_resolve ( eth_iface_t * iface,  
                  ipv4_addr_t dest,  
                  mac_addr_t mac );
```

